



## 0. Stem network



```
# stem net
self.conv1 = nn.Conv2d(3, 64, kernel_size=3, stride=2, padding=1, bias=False)
self.bn1 = self.norm_layer(64)
self.conv2 = nn.Conv2d(64, 64, kernel_size=3, stride=2, padding=1, bias=False)
self.bn2 = self.norm_layer(64)
self.relu = nn.ReLU(inplace=True)
```

Stride=2의 convolution을 2번 거쳐서 h/4 x w/4 크기로 만들

## Main body

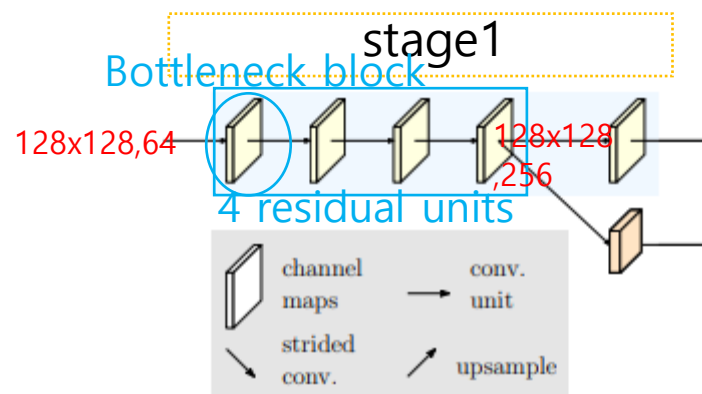


### 3.4 Instantiation

The main body contains four stages with four parallel convolution streams. The resolutions are  $1/4$ ,  $1/8$ ,  $1/16$ , and  $1/32$ . The first stage contains 4 residual units where each unit is formed by a bottleneck with the width 64, and is followed by one  $3 \times 3$  convolution changing the width of feature maps to  $C$ . The 2<sup>nd</sup>, 3<sup>rd</sup>, 4<sup>th</sup> stages contain 1, 4, 3 modularized blocks, respectively. Each branch in multi-resolution parallel convolution of the modularized block contains 4 residual units. Each unit contains two  $3 \times 3$  convolutions for each resolution, where each convolution is followed by batch normalization and the nonlinear activation ReLU. The widths (numbers of channels) of the

Hrnet32  
 $C = 32$

# 1. Main Body – stage1



```
HRNET_32.STAGE1 = CN()
HRNET_32.STAGE1.NUM_MODULES = 1
HRNET_32.STAGE1.NUM_BRANCHES = 1
HRNET_32.STAGE1.NUM_BLOCKS = [4] # first stage contains 4 residual unit
HRNET_32.STAGE1.BLOCK = 'BOTTLENECK' # where each unit is formed by a bottleneck
HRNET_32.STAGE1.NUM_CHANNELS = [64] # with width 64
HRNET_32.STAGE1.FUSE_METHOD = 'SUM'
```

Table 14

Resolution	Stage 1	Stage 2	Stage 3	Stage 4
4×	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 4 \times 1$	$\begin{bmatrix} 3 \times 3, C \\ 3 \times 3, C \end{bmatrix} \times 4 \times 1$	$\begin{bmatrix} 3 \times 3, C \\ 3 \times 3, C \end{bmatrix} \times 4 \times 4$	$\begin{bmatrix} 3 \times 3, C \\ 3 \times 3, C \end{bmatrix} \times 4 \times 3$
8×		$\begin{bmatrix} 3 \times 3, 2C \\ 3 \times 3, 2C \end{bmatrix} \times 4 \times 1$	$\begin{bmatrix} 3 \times 3, 2C \\ 3 \times 3, 2C \end{bmatrix} \times 4 \times 4$	$\begin{bmatrix} 3 \times 3, 2C \\ 3 \times 3, 2C \end{bmatrix} \times 4 \times 3$
16×			$\begin{bmatrix} 3 \times 3, 4C \\ 3 \times 3, 4C \end{bmatrix} \times 4 \times 4$	$\begin{bmatrix} 3 \times 3, 4C \\ 3 \times 3, 4C \end{bmatrix} \times 4 \times 3$
32×				$\begin{bmatrix} 3 \times 3, 8C \\ 3 \times 3, 8C \end{bmatrix} \times 4 \times 3$

# 1. Main Body – stage1




Table 14

Resolution	Stage 1	
4x	<div><div><div>1 × 1, 64</div><div>3 × 3, 64</div><div>1 × 1, 256</div></div><div>× 4 × 1</div></div>	<div><div>3 ×</div><div>3 ×</div></div>
8x	Bottleneck block	<div><div>3 ×</div><div>3 ×</div></div>
16x		
32x		

class Bottleneck(nn.Module):

expansion = 4

```
def __init__(self, inplanes, planes, stride=1, downsample=None, groups=1,
              base_width=64, dilation=1, norm_layer=None):
    super(Bottleneck, self).__init__()
    if norm_layer is None:
        norm_layer = nn.BatchNorm2d
    width = int(planes * (base_width / 64.)) * groups
    # Both self.conv2 and self.downsample layers downsample the input when stride != 1
    self.conv1 = conv1x1(inplanes, width)
    self.bn1 = norm_layer(width)
    self.conv2 = conv3x3(width, width, stride, groups, dilation)
    self.bn2 = norm_layer(width)
    self.conv3 = conv1x1(width, planes * self.expansion)
    self.bn3 = norm_layer(planes * self.expansion)
    self.relu = nn.ReLU(inplace=True)
    self.downsample = downsample
    self.stride = stride
```

```
def forward(self, x):
    identity = x

    out = self.conv1(x)
    out = self.bn1(out)
    out = self.relu(out)

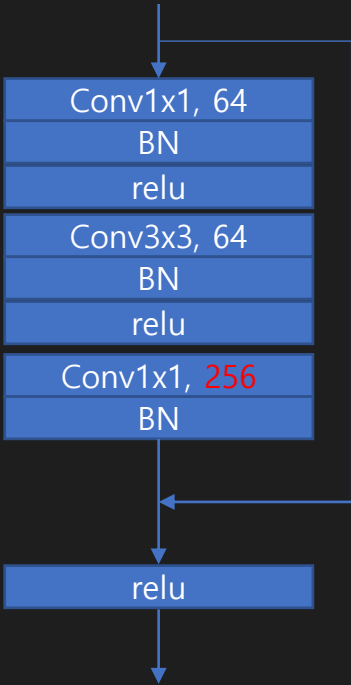
    out = self.conv2(out)
    out = self.bn2(out)
    out = self.relu(out)

    out = self.conv3(out)
    out = self.bn3(out)


    if self.downsample is not None:
        identity = self.downsample(x)

    out += identity
    out = self.relu(out)

    return out
```



# 1. Main Body – stage1



```
def _make_layer(self, block, inplanes, planes, blocks, stride=1):
    downsample = None
    if stride != 1 or inplanes != planes * block.expansion:
        downsample = nn.Sequential(
            nn.Conv2d(inplanes, planes * block.expansion,
                      kernel_size=1, stride=stride, bias=False),
            self.norm_layer(planes * block.expansion),
        )

    layers = []
    layers.append(block(inplanes, planes, stride, downsample, norm_layer=self.norm_layer))

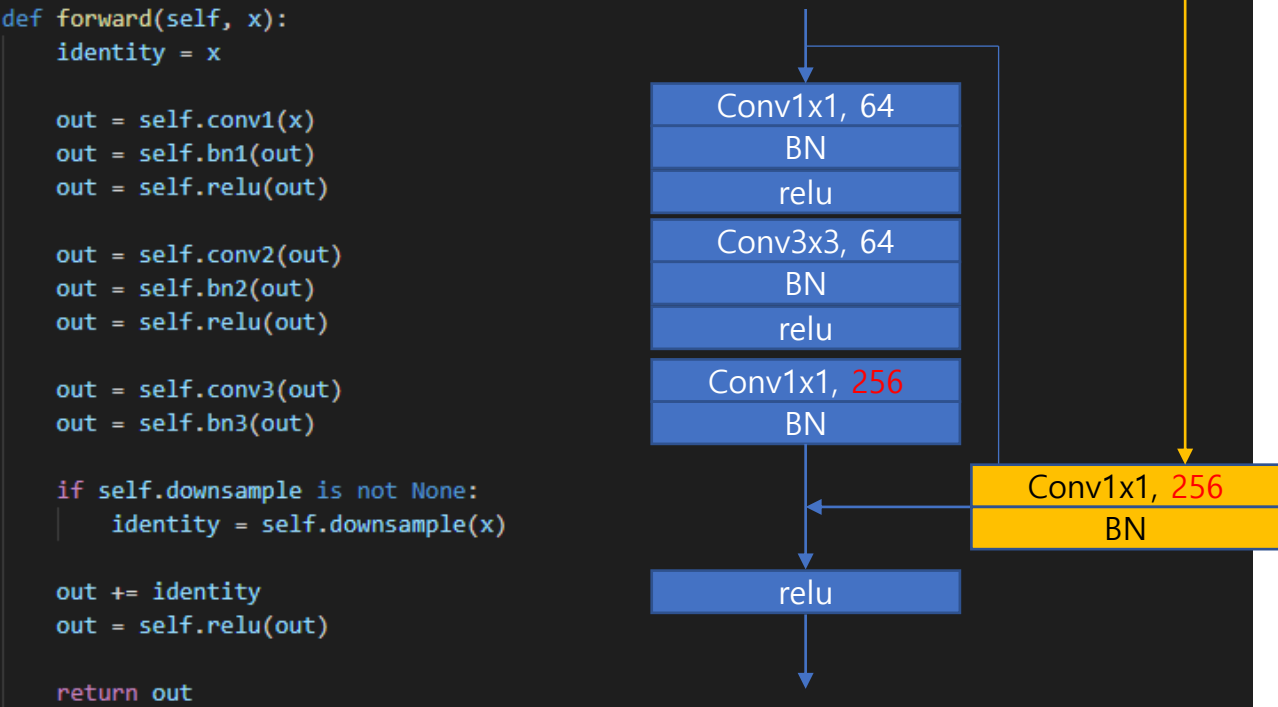
    inplanes = planes * block.expansion
    for _ in range(1, blocks): # 나머지 3개 unit마다 class Bottleneck(nn.Module) 객체를 만들어서 이어줌
        layers.append(block(inplanes, planes, norm_layer=self.norm_layer))

    return nn.Sequential(*layers)
```


첫 번째 block은 입력 채널(64)과 출력채널(256)을 맞추기 위해 conv1x1 추가

Table 14

Resolution	Stage 1	
4x	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 4 \times 1$	$\begin{bmatrix} 3 \times \\ 3 \times \end{bmatrix}$
8x		$\begin{bmatrix} 3 \times \\ 3 \times \end{bmatrix}$
16x		
32x		




## 1. Main Body – stage1



# 1. Main Body – stage1

$$\begin{array}{ccccccc}
 \mathcal{N}_{11} & \rightarrow & \mathcal{N}_{21} & \rightarrow & \mathcal{N}_{31} & \rightarrow & \mathcal{N}_{41} \\
 & \searrow & & & & & \\
 & & \mathcal{N}_{22} & \rightarrow & \mathcal{N}_{32} & \rightarrow & \mathcal{N}_{42} \\
 & & & \searrow & & & \\
 & & & & \mathcal{N}_{33} & \rightarrow & \mathcal{N}_{43} \\
 & & & & & \searrow & \\
 & & & & & & \mathcal{N}_{44},
 \end{array}$$



```

def _make_transition_layer(self, num_channels_pre_layer, num_channels_cur_layer):

    num_branches_cur = len(num_channels_cur_layer) # num_channels_cur_layer : [32, 64]
    num_branches_pre = len(num_channels_pre_layer) # num_channels_pre_layer : [64]

    transition_layers = []
    for i in range(num_branches_cur):
        if i < num_branches_pre:
            if num_channels_cur_layer[i] != num_channels_pre_layer[i]:
                transition_layers.append(nn.Sequential(
                    nn.Conv2d(num_channels_pre_layer[i],
                              num_channels_cur_layer[i],
                              3,
                              1,
                              1,
                              bias=False),
                    self.norm_layer(num_channels_cur_layer[i]),
                    nn.ReLU(inplace=True)))
            else:
                transition_layers.append(None)
        else:
            conv3x3s = []
            for j in range(i+1-num_branches_pre):
                inchannels = num_channels_pre_layer[-1]
                outchannels = num_channels_cur_layer[i] \
                    if j == i-num_branches_pre else inchannels
                conv3x3s.append(nn.Sequential(
                    nn.Conv2d(
                        inchannels, outchannels, 3, 2, 1, bias=False),
                    self.norm_layer(outchannels),
                    nn.ReLU(inplace=True)))
            transition_layers.append(nn.Sequential(*conv3x3s))

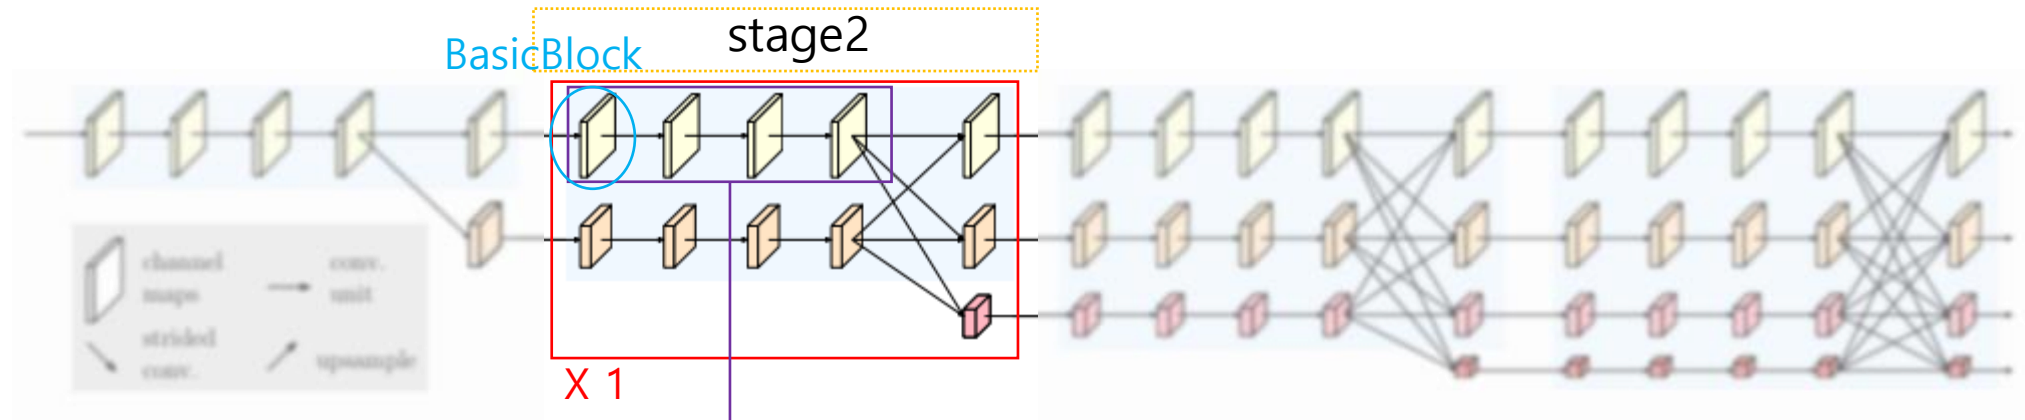
    return nn.ModuleList(transition_layers)
    
```

Resolution 그대로 & Channel 수 변환

Resolution 그대로 & Channel 수 그대로

Resolution 절반 & Channel 수 두 배

# 1. Main Body – stage2



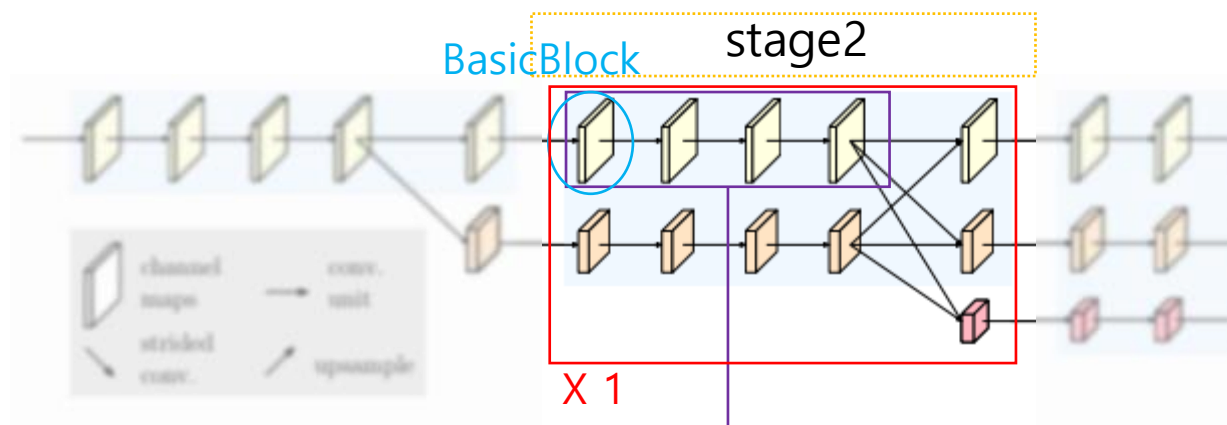
Each branch contains 4 residual units

2<sup>nd</sup> stage contain  
1 modularized blocks

Resolution	Stage 1	Stage 2	Stage 3	Stage 4
4×	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 4 \times 1$	$\begin{bmatrix} 3 \times 3, C \\ 3 \times 3, C \end{bmatrix} \times 4 \times 1$	$\begin{bmatrix} 3 \times 3, C \\ 3 \times 3, C \end{bmatrix} \times 4 \times 4$	$\begin{bmatrix} 3 \times 3, C \\ 3 \times 3, C \end{bmatrix} \times 4 \times 3$
8×		$\begin{bmatrix} 3 \times 3, 2C \\ 3 \times 3, 2C \end{bmatrix} \times 4 \times 1$	$\begin{bmatrix} 3 \times 3, 2C \\ 3 \times 3, 2C \end{bmatrix} \times 4 \times 4$	$\begin{bmatrix} 3 \times 3, 2C \\ 3 \times 3, 2C \end{bmatrix} \times 4 \times 3$
16×			$\begin{bmatrix} 3 \times 3, 4C \\ 3 \times 3, 4C \end{bmatrix} \times 4 \times 4$	$\begin{bmatrix} 3 \times 3, 4C \\ 3 \times 3, 4C \end{bmatrix} \times 4 \times 3$
32×				$\begin{bmatrix} 3 \times 3, 8C \\ 3 \times 3, 8C \end{bmatrix} \times 4 \times 3$



# 1. Main Body – stage2



Each branch contains 4 residual units

2<sup>nd</sup> stage contain  
1 modularized blocks

Resolution	Stage 1	Stage 2
4×	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 4 \times 1$	$\begin{bmatrix} 3 \times 3, C \\ 3 \times 3, C \end{bmatrix} \times 4 \times 1$
8×		$\begin{bmatrix} 3 \times 3, 2C \\ 3 \times 3, 2C \end{bmatrix} \times 4 \times 1$
16×		
32×		

```
def _make_stage(self, layer_config, num_inchannels,
                multi_scale_output=True):
    num_modules = layer_config['NUM_MODULES']
    num_branches = layer_config['NUM_BRANCHES']
    num_blocks = layer_config['NUM_BLOCKS']
    num_channels = layer_config['NUM_CHANNELS']
    block = blocks_dict[layer_config['BLOCK']]
    fuse_method = layer_config['FUSE_METHOD']


    modules = []
    for i in range(num_modules):
        # multi_scale_output is only used last module
        if not multi_scale_output and i == num_modules - 1:
            reset_multi_scale_output = False
        else:
            reset_multi_scale_output = True

        modules.append(
            HighResolutionModule(num_branches,
                                block,
                                num_blocks,
                                num_inchannels,
                                num_channels,
                                fuse_method,
                                reset_multi_scale_output,
                                norm_layer=self.norm_layer)
        )
    num_inchannels = modules[-1].get_num_inchannels()

    return nn.Sequential(*modules), num_inchannels
```

Stage2: 17  
Stage3: 47  
Stage4: 37


# 1. Main Body – stage2



stage2에서는 2개,  
stage3에서는 3개,  
Stage4에서는 4개의 branch를 더함

```
class HighResolutionModule(nn.Module):  
    ...  
    def _make_branches(self, num_branches, block, num_blocks, num_channels):  
        branches = []  
  
        for i in range(num_branches):  
            branches.append(  
                self._make_one_branch(i, block, num_blocks, num_channels))  
  
        return nn.ModuleList(branches)
```

# 1. Main Body – stage2



Each branch contains 4 residual units

2<sup>nd</sup> stage contain  
1 modularized blocks


stage2에서는 2개,  
stage3에서는 3개,  
Stage4에서는 4개의  
branch(BasicBlock)을 더해줌

```
class HighResolutionModule(nn.Module):  Modularized block
    ...
    def _make_branches(self, num_branches, block, num_blocks, num_channels):
        branches = []

        for i in range(num_branches):
            branches.append(
                self._make_one_branch(i, block, num_blocks, num_channels))

        return nn.ModuleList(branches)
```

## 1. Main Body – stage2



```
class HighResolutionModule(nn.Module):
    ...
    def _make_one_branch(self, branch_index, block, num_blocks, num_channels,
                        stride=1):
        downsample = None
        if stride != 1 or \
            self.num_inchannels[branch_index] != num_channels[branch_index] * block.expansion:
            downsample = nn.Sequential(
                nn.Conv2d(self.num_inchannels[branch_index],
                        num_channels[branch_index] * block.expansion,
                        kernel_size=1, stride=stride, bias=False),
                self.norm_layer(num_channels[branch_index] * block.expansion),
            )

        layers = []
        layers.append(block(self.num_inchannels[branch_index],
                            num_channels[branch_index], stride, downsample, norm_layer=self.norm_layer))
        self.num_inchannels[branch_index] = \
            num_channels[branch_index] * block.expansion
        for i in range(1, num_blocks[branch_index]):
            layers.append(block(self.num_inchannels[branch_index],
                                num_channels[branch_index], norm_layer=self.norm_layer))


        return nn.Sequential(*layers)
```

branch 하나

- residual unit (BasicBlock) 4개로 이루어짐
- Each branch in multi-resolution parallel convolution of the modularized block contains **4 residual units**.

# 1. Main Body – stage2

## residual unit (BasicBlock)



```
class BasicBlock(nn.Module):  
    expansion = 1  
  
    def __init__(self, inplanes, planes, stride=1, downsample=None, groups=1,  
                 base_width=64, dilation=1, norm_layer=None):  
        super(BasicBlock, self).__init__()  
        if norm_layer is None:  
            norm_layer = nn.BatchNorm2d  
        if groups != 1 or base_width != 64:  
            raise ValueError('BasicBlock only supports groups=1 and base_width=64')  
        if dilation > 1:  
            raise NotImplementedError("Dilation > 1 not supported in BasicBlock")  
        # Both self.conv1 and self.downsample layers downsample the input when stride != 1  
        self.conv1 = conv3x3(inplanes, planes, stride)  
        self.bn1 = norm_layer(planes)  
        self.relu = nn.ReLU(inplace=True)  
        self.conv2 = conv3x3(planes, planes)  
        self.bn2 = norm_layer(planes)  
        self.downsample = downsample  
        self.stride = stride  
  
    def forward(self, x):  
        identity = x  
  
        out = self.conv1(x) # 3x3  
        out = self.bn1(out)  
        out = self.relu(out)  
  
        out = self.conv2(out) # 3x3  
        out = self.bn2(out)  
  
        if self.downsample is not None:  
            identity = self.downsample(x)  
  
        out += identity  
        out = self.relu(out)  
  
        return out
```

```
graph TD  
    x((x)) --> J1(( ))  
    J1 --> C1[Conv3x3]  
    C1 --> B1[BN]  
    B1 --> R1[relu]  
    R1 --> J2(( ))  
    J1 --> C2[Conv3x3]  
    C2 --> B2[BN]  
    B2 --> C3[Conv1x1]  
    C3 --> B3[BN]  
    B3 --> J2  
    J2 --> R2[relu]  
    R2 --> out((out))
```

## 1. Main Body – stage2





Fig. 3. Illustrating how the fusion module aggregates the information for high, medium and low resolutions from left to right, respectively. Right legend: strided  $3 \times 3$  = stride-2  $3 \times 3$  convolution, up samp.  $1 \times 1$  = bilinear upsampling followed by a  $1 \times 1$  convolution.

- Strided Convolution으로 하위 stream 생성
- Bilinear upsampling  $\rightarrow$   $1 \times 1$  Convolution으로 상위 stream 생성

# 1. Main Body – stage2

```
def _make_fuse_layers(self): # high resolution으로 만들때는 bilinear->conv1x1, low resolution으로 만들때는 conv3x3 (stride=2)
    if self.num_branches == 1:
        return None

    num_branches = self.num_branches
    num_inchannels = self.num_inchannels
    fuse_layers = []
    for i in range(num_branches if self.multi_scale_output else 1):
        fuse_layer = []
        for j in range(num_branches):
            if j > i: # up samp. 1 x 1 = bilinear upsampling followed by a 1 x 1 convolution.
                fuse_layer.append(nn.Sequential(
                    nn.Conv2d(num_inchannels[j],
                              num_inchannels[i],
                              1,
                              1,
                              0,
                              bias=False),
                    self.norm_layer(num_inchannels[i])))
            elif j == i:
                fuse_layer.append(None)
            else: # strided 3 x 3 = stride-2 3 x 3 convolution
                conv3x3s = []
                for k in range(i-j):
                    if k == i - j - 1:
                        num_outchannels_conv3x3 = num_inchannels[i]
                        conv3x3s.append(nn.Sequential(
                            nn.Conv2d(num_inchannels[j],
                                        num_outchannels_conv3x3,
                                        3, 2, 1, bias=False),
                            self.norm_layer(num_outchannels_conv3x3)))
                    else:
                        num_outchannels_conv3x3 = num_inchannels[j]
                        conv3x3s.append(nn.Sequential(
                            nn.Conv2d(num_inchannels[j],
                                        num_outchannels_conv3x3,
                                        3, 2, 1, bias=False),
                            self.norm_layer(num_outchannels_conv3x3),
                            nn.ReLU(inplace=True)))
                fuse_layer.append(nn.Sequential(*conv3x3s))
        fuse_layers.append(nn.ModuleList(fuse_layer))

    return nn.ModuleList(fuse_layers)
```





Fig. 3. Illustrating how the fusion module aggregates the information for high, medium and low resolutions from left to right, respectively. Right legend: strided  $3 \times 3$  = stride-2  $3 \times 3$  convolution, up samp.  $1 \times 1$  = bilinear upsampling followed by a  $1 \times 1$  convolution.

# 1. Main Body – stage2



```
class HighResolutionModule(nn.Module): Modularized block
    ...
    def forward(self, x):
        if self.num_branches == 1:
            return [self.branches[0](x[0])]

        for i in range(self.num_branches):
            x[i] = self.branches[i](x[i])

        x_fuse = []
        for i in range(len(self.fuse_layers)):
            y = x[0] if i == 0 else self.fuse_layers[i][0](x[0])
            for j in range(1, self.num_branches):
                if i == j:
                    y = y + x[j]
                elif j > i:
                    width_output = x[j].shape[-1]
                    height_output = x[j].shape[-2]
                    y = y + F.interpolate(
                        self.fuse_layers[i][j](x[j]),
                        size=[height_output, width_output],
                        mode='bilinear',
                        align_corners=True
                    )
                else:
                    y = y + self.fuse_layers[i][j](x[j])
            x_fuse.append(self.relu(y))

        return x_fuse
```



## 2. Last layer

```
# Upsampling
x0_h, x0_w = x[0].size(2), x[0].size(3)
x1 = F.interpolate(x[1], size=(x0_h, x0_w), mode='bilinear', align_corners=False)
x2 = F.interpolate(x[2], size=(x0_h, x0_w), mode='bilinear', align_corners=False)
x3 = F.interpolate(x[3], size=(x0_h, x0_w), mode='bilinear', align_corners=False)

x = torch.cat([x[0], x1, x2, x3], 1)

x = self.last_layer(x) # conv1x1(c->c) -> bn -> relu -> conv1x1(c->19)

return x
```

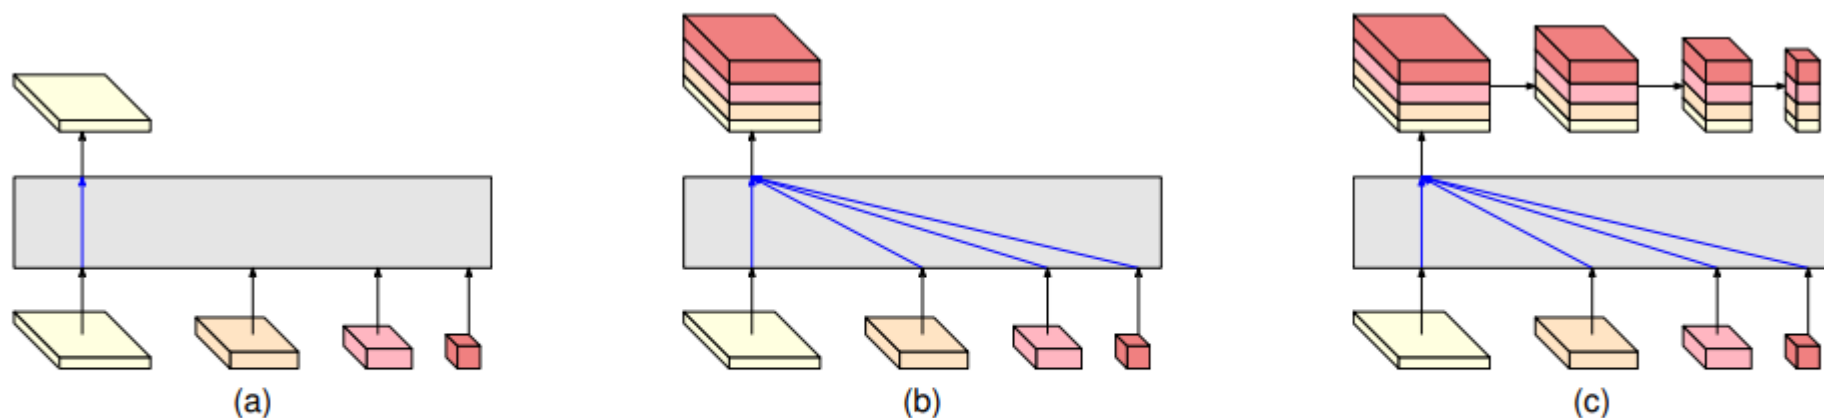


Fig. 4. (a) HRNetV1: only output the representation from the high-resolution convolution stream. (b) HRNetV2: Concatenate the (upsampled) representations that are from all the resolutions (the subsequent  $1 \times 1$  convolution is not shown for clarity). (c) HRNetV2p: form a feature pyramid from the representation by HRNetV2. The four-resolution representations at the bottom in each sub-figure are outputted from the network in Figure 2, and the gray box indicates how the output representation is obtained from the input four-resolution representations.